

Packaging and Deployment

- Java Packaging for Bare Metal
 - JAR Files. Fat (Uber) JARs
 - JPMS. `jlink/jpackage`
 - GraalVM `native-image`
- Linux Packaging for Bare Metal (DEB, RPM)
- Virtualization and Containers 101. Containers: Concept, Implementation, Solutions. Container Orchestration
- Containers and Java: Build. Manual Dockerfile, `docker-maven-plugin`, Google Jib
- Containers and Java: Runtime

JAR (Java ARchive)

<https://docs.oracle.com/en/java/javase/17/docs/specs/jar/jar.html>

- **ZIP archive** with **compiled Java classes**, **resources** and **metainformation** (META-INF/)
- **Compiled classes** and **class resources** are put into directories corresponding to Java packages.
 - Top-Level class `ru.hse.java.HelloWorld` => `ru/hse/java/HelloWorld.class`
 - Anonymous, inner, static inner: `ru.hse.HelloWorld.Inner` => `ru/hse/HelloWorld$Inner.class`
- Most important **metainformation** is the Manifest, META-INF/MANIFEST.MF:

Manifest-Version: 1.0

Main-Class: `ru.hse.MyCoolProjectMain`
<more key-value pairs...>

`java -jar my-cool-project-1.0.jar`

- META-INF/ directory also MAY contain:
 - Digital signature files (*.RSA, *.DSA, SIG-*)
 - **Service Providers** (META-INF/services/<fully-qualified name of Service Class Impl>)
@see future seminar on DI

Fat (Uber) JARs

<https://stackoverflow.com/a/36539885/3438672>

- Bring all of your dependencies (transitively) into a single large JAR file
- Three methods:
 - **Unpack all dependency JARs** (`maven-assembly-plugin` → `jar-with-dependencies`)
 - Unpack all dependency JARs, but also rename their packages and merge resources (`maven-shade-plugin`)
 - To avoid conflicts with your library users' dependencies
 - Copy JARs into your JAR and use a special *class loader* to transparently access their classes (`onejar`, `spring-boot-plugin:repackage`)
- Classpath is Linear. Multiple Versions of Artifact on Class Path → JAR Hell
 - `maven-enforcer-plugin` helps by forcing **dependency convergence**. Use it!
<https://maven.apache.org/enforcer/enforcer-rules/dependencyConvergence.html>
 - Explicitly divide your code into modules: JPMS (Java 9+ standard), OSGi (more flexible)

JPMS. jlink/jpackage

<https://nipafx.dev/java-module-system-tutorial/>

– Java Platform Module System

- Dependencies (**requires** [transitive] ...). Cycles and Split Packages are disallowed
- Strong Encapsulation. Visibility: **exports** x [to z], Reflection Access (**opens** y [to x])
- Inter-module dependencies are assumed to be *mostly* static (determined at app start time)
- Requires you to **cleanly separate your system into modules**, which is **HARD**
- jlink image will include only modules used by your app, including JDK modules:

```
jlink --module-path $JAVA_HOME/jmods:mlib # JDK \  
      --add-modules com.greetings          # Your Main Module \  
      --output greetingsapp
```

<https://dev.java/learn/creating-runtime-and-application-images-with-jlink/>

- jpackage can create native packages for Windows (.msi), Mac OS (.dmg) and Linux
- **Non-modular apps are supported, too!**

<https://dev.java/learn/jpackage/#examples>

<https://medium.com/azulsystems/9568c5e70ef4>

GraalVM native-image

- Unlike `jlink` which only links **platform-independent** modules together, native-image Ahead-of-Time compiles your code into a **platform-dependent native application**
 - Some limitations, mostly around Reflection usage and static class initialization (**static finals** and **static {}** blocks spawning threads and the like)
 - Popular frameworks are supported out-of-box, less popular can encounter problems
- Emerging frameworks using native-image as the default: Quarkus, Micronaut
- Good fit for Microservices, Serverless, CLIs, Agents, ...
 - **Fast start** (low start latency) far more important than **peak throughput**
 - GC and runtime is naive compared to OpenJDK's Hotspot

Linux Packaging (DEB, RPM)

- General tools for managing components of a Linux distribution
 - Manage *packages*, which provide *files* installed into your filesystem, optionally with *hook* scripts executed on package installation and removal
 - Packages are *versioned* and *depend* on each other (specific version or version range)
 - High-Level Package Managers (e.g. apt, yum, zypper): Dependency Resolution, Repository Management
 - Low-Level Package Managers (e.g. dpkg, rpm): Local package DB, Package installation
 - Package authoring tools (e.g. rpmbuild): Transform specs, source files and scripts into a single coherent package
- **Cannot have multiple versions of the same package installed at the same time!**
 - Multiple packages, however, *can* provide the same binary
E.g. /bin/java via update-java-alternatives (basically, symlinks)

Virtualization

Run 1+ *Guests* (OS w/virtualized hardware) on a single *Host*, managed by a *Hypervisor*
Reasonably fast: hardware (e.g., Intel VT-x) and software-assisted (*Paravirtualization*)

Advantages (over bare metal)	Drawbacks
<ul style="list-style-type: none">• Tight Isolation (CPU, RAM, Network, Storage, ...)• Better Resource Utilization: A single host server can manage multiple guest Vms• Better Scalability: both Vertical (allocate more host resources to guest) and Horizontal (spawn more guests)• Better Resilience and Disaster Recovery:<ul style="list-style-type: none">• VM crashed? Just spawn a new one instead• Replicated network block storage	<ul style="list-style-type: none">• You bring <i>Everything but the Kitchen Sink</i>: Full OS image + all packages + your app (orders of magnitude smaller)• Expensive VM setup and teardown: Tens of seconds to minutes, depending on workload and resources available)• Some performance degradation (top → Steal time) Especially in burstable/preemptible VMs


Containers

- Hypervisors virtualize hardware and then run a full-fledged real OS on it
- But modern OSs **already have powerful primitives** for isolation (process, network, storage) and resource constraints!

Container includes all relevant (software and data, but **NOT** OS kernel/image) dependencies for your application/service

...unlike VMs, you cannot *e.g.* run it on a fixed specific kernel version.
But this is rarely required for most common apps and services anyway

Containers: Packaging & Delivery

- Linux Containers were mostly intended for **standard packaging and delivery** of software
 - Compare e.g. FreeBSD Jails, which were designed **as a security feature** first and foremost
- Docker Images → OCI (container image standard)
 - Optimized for **efficient download and caching**:
 - Layered filesystem
 - Hash-based layer identity (SHA256 digest)
 - Images are **tagged**, typically with version. latest tag (convention)
- Container Registries are somewhat similar to Maven artifact repositories, but **content-addressable** via layer hashes
 - Public, e.g. DockerHub
 - Private. Cloud offerings from all major cloud providers, including **Yandex**  **Cloud**

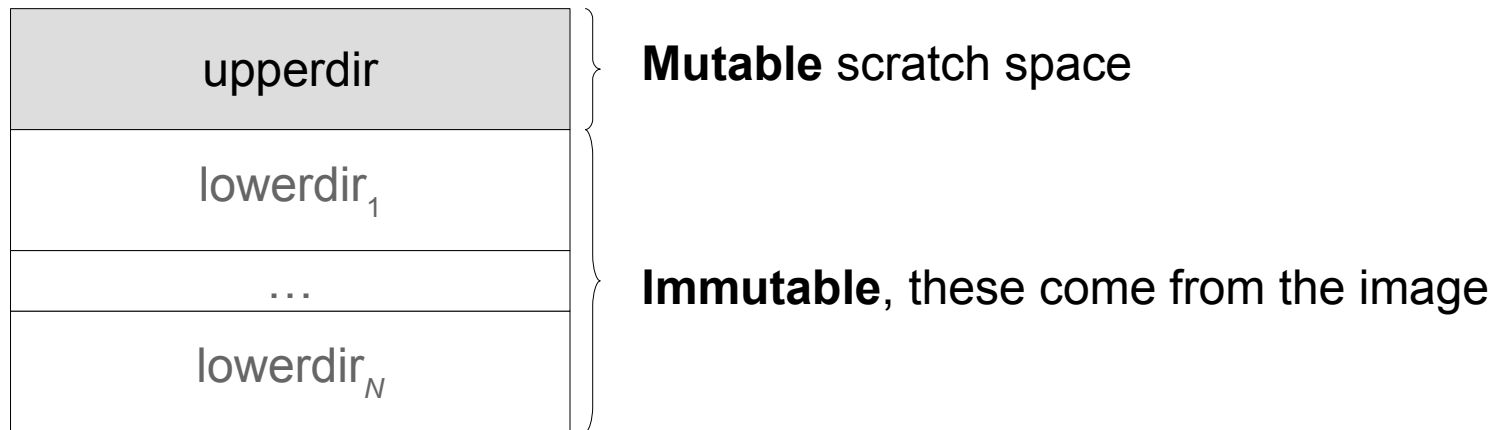
Containers: Isolation → Namespaces

Unit of isolation in Containers is the **Process**

- Namespaces limit **visibility** of sensitive system entities: processes, network interfaces, users, mount points, ...
- **Process isolation**: main process in container runs as PID 1, but its PID in the host is different
- **Network isolation**: container has limited access to host network interfaces, can have special *bridge* network interfaces etc.
- **User isolation**: process inside container runs as root (or as some user)
- **Filesystem isolation**: containers cannot access host storage unless explicitly specified

Containers: Isolation → FS

- `pivot_root` (chroot on steroids) to have / independent of host root filesystem
- Union Filesystem, *overlayfs*



Containers: Limits and Security

- cgroups (Control Groups): Hierarchical Resource Accounting and Limits
 - Limit CPU core usage
 - Memory usage (RSS, resident set size)
 - I/O (read and write iops)
 - Network bandwidth
 - Process is typically killed or throttled when resource overuse is detected
- Container Security: Principle of least privilege
 - Drop capabilities (CAP_...)
 - seccomp-bpf: Selective filtering of syscalls

Container Orchestration

- Containers are much more lightweight than VMs
 - You can have tens and even hundreds of them running on a modest VM!
- Container Orchestration (*E.g.* Kubernetes, Docker Swarm, Nomad, Amazon ECS, ...): Managing lots of containers, and Clusters of similar containers
 - Resource Allocator
 - Workload Scheduler
 - Jobs: batch, throughput-oriented workloads
 - Services (stateless and stateful): interactive, latency-oriented workloads
 - Resilience: restarts/retries, container migration
 - Autoscaling
 - Persistent storage management
- Container Management Philosophy: Cattle **NOT** Pets!

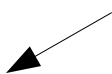
Popular Container Solutions

- **Basic Container Management: *Docker*** (safer, simpler production alternatives are *containerd* and *rkt*)
 - build image, push (publish to registry), pull (download from registry)
 - run image/run command inside image
 - view stdout/stderr of container process/route process logs to syslog and whatever
- **Basic Orchestration: *Docker Compose***
- **Advanced Orchestration: *Kubernetes* / *OpenShift*, *Amazon ECS*, *Mesosphere DC/OS*, *Hashicorp Nomad***
- **Container-Optimized Linux distributions (*Alpine*, *distroless*):**
Minimal dependencies required to run containers →
→ Smaller attack surface and better performance
- **Container-First Operating Systems: *CoreOS*, *AWS Bottlerocket***

Containers and Java: Manual Build

```
src/main/docker/Dockerfile
```

```
FROM bellsoft/liberica-openjdk-alpine-musl:17.0.3-7
RUN apk add strace
COPY echo-server-1.0-SNAPSHOT-jar-with-dependencies.jar /
CMD ["java", "-jar", "/echo-server-1.0-SNAPSHOT-jar-with-dependencies.jar"]
EXPOSE 48484
```



Fat JAR!

Terminal

```
$ mvn clean package && docker build target/ \ # Build directory
-t cr.yandex/crp024d3b67qde0rn13r/echo-server:latest \ # Tag
-f src/main/docker/Dockerfile # Path to Dockerfile
```

Containers and Java: Manual Run

Terminal 1

Ports

Docker Image

```
$ docker run --rm -it -p 48484:48484 cr.yandex/crp024d3b67qde0rn13r/echo-server:latest
```

Run Interactive, with Terminal attached

Cleanup on Termination

Terminal 2

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND NAMES	CREATED
f53c9e43ab97	cr.yandex/crp024d3b67qde0rn13r/echo-server:latest	"java -jar /ech..."	3 seconds ago
Up 1 second	0.0.0.0:48484->48484/tcp, :::48484->48484/tcp	recursing_hawking	

```
$ telnet localhost 48484
```

```
Trying 127.0.0.1...
```

```
Connected to localhost.
```

```
Escape character is '^]'.  
Hi
```

```
Hi
```

```
END
```

```
Connection closed by foreign host.
```


docker-maven-plugin

<https://github.com/fabric8io/docker-maven-plugin>

- Modeled after maven-assembly-plugin, but dependencies are packaged into a Docker container instead of a Fat JAR
- Dockerfile map to <tags> in pom.xml:

```
<from>bellsoft/liberica-openjdk-alpine-musl:17.0.3-7</from>  
<assembly><descriptorRef>artifact-with-dependencies</descriptorRef></assembly>  
<cmd><exec>  
    <arg>java</arg><arg>-jar</arg><arg>maven/${project.build.finalName}.jar</arg>  
</exec></cmd>
```
- Needs docker binaries for the build (Mac users beware!)
- After build, you have to manually tag built image and push it:

```
$ docker tag echo-server:1.0-SNAPSHOT cr.yandex/crp024d3b67qde0rn13r/echo-server:latest  
$ docker push cr.yandex/crp024d3b67qde0rn13r/echo-server:latest
```

Google Jib (Java Image Builder)

<https://github.com/GoogleContainerTools/jib>

- Builds an **Efficient Layered Image**: your code in one layer, dependencies in the other
- Does not need docker binaries to run
- Pushes the built image to your Docker repository
- Convention over Configuration:

```
<from>
```

```
    <image>bellsoft/libERICA-openjdk-alpine-musl:17.0.3-7</image>
```

```
</from>
```

} Optional!

```
<to>
```

```
    <image>cr.yandex/<...>/${project.artifactId}:latest</image>
```

```
</to>
```

Tutorial: <https://phauer.com/2019/no-fat-jar-in-docker-image/>

Containers and Java: Runtime

- Enable namespace and cgroup-related java command line switches, e.g.:
 - XX:+UseContainerSupport -XX:InitialRAMPercentage=80 -XX:MaxRAMPercentage=80
- Getting memory and CPU limits is **REALLY HARD** and needs a lot of fine tuning, even for experts. @see <https://www.youtube.com/watch?v=kKigibHrV5I>
- To properly catch SIGTERM, use explicit argv[] in Dockerfile:

```
CMD ["java", "-jar", "/your-app.jar"]  
# This runs shell, do NOT use: CMD java -jar /your-app.jar
```
- A shell script which runs java? It **MUST** properly handle SIGTERM!
 - Running java with exec is the simplest solution
 - In more complex cases, you might need to register a trap handler in the script
- You run as PID 1 and need to take care of Zombie processes. `docker run --init`
- App stderr and stdout: `docker logs`, @see <https://docs.docker.com/config/containers/logging/>